



TECHNOLOGICAL INSTITUTE OF THE PHILIPPINES
938 Aurora Blvd., Cubao, Quezon City

COLLEGE OF ENGINEERING AND ARCHITECTURE
ELECTRONICS ENGINEERING DEPARTMENT

1ST SEMESTER SY 2022 - 2023

Prediction and Machine Learning

COE 005
ECE41S11

Generative Adversarial Network

Generate Cartoon Characters

Submitted to:
Engr. Christian Lian Paulo Rioflorido

Submitted on:
October 21, 2022

Submitted by:
So, Don Gabriel L.

INTRODUCTION:

GAN is a deep learning technique used to generate some new data from scratch. It also flourishes the what so-called “This waifu does not exist” which becomes a popular topic amongst anime fans due to how you can generate your own “waifu” with just a click. GAN has two networks, the generator and the discriminator network wherein the generator will generate some fake data while the discriminator will be the one to identify a couple of data in the generated fake data as well as the real data. The generator will mimic the real data whereas the discriminator will train its neural network for detecting the fake and real data, both networks will try to improve one another's performance through backpropagation.

```
10000 [Discriminator loss: 0.481456, acc.: 76.31%] [Generator loss: 2.071867]  
1/1 [=====] - 0s 20ms/step
```



Generated Anime Face at 10k epochs

PROCEDURE AND SCREENSHOTS OF THE PROGRAM:

The editor used for this exercise is the Google Collab.

1. First, import the libraries needed for the program.

```
import numpy as np
import os
import matplotlib.pyplot as plt
import cv2

import warnings
warnings.filterwarnings('ignore')

import keras
from tensorflow.keras.optimizers import Adam
from keras.models import Sequential, Model
from keras.layers import Dense, LeakyReLU, Reshape, Flatten, Input
from keras.layers import Conv2D, MaxPooling2D, Activation, Dropout, Conv2DTranspose

from tensorflow.compat.v1.keras.layers import BatchNormalization
```

2. This is followed by the loading of the data as well as data processing. Here, we created a function where only the set file type (.jpg, .jpeg, .png, .bmp) will be taken and used by the program. We also downloaded the dataset through the code, you just need to have the username and key of your Kaggle account since the url of our dataset is from Kaggle and after downloading, we will store in a variable, which is named images in this program.

```
[ ] def list_images(basePath, contains=None):
    # return the set of files that are valid
    return list_files(basePath, validExts=(".jpg", ".jpeg", ".png", ".bmp"), contains=contains)

def list_files(basePath, validExts=(".jpg", ".jpeg", ".png", ".bmp"), contains=None):
    # loop over the directory structure
    for (rootDir, dirNames, filenames) in os.walk(basePath):
        # loop over the filenames in the current directory
        for filename in filenames:
            # if the contains string is not none and the filename does not contain
            # the supplied string, then ignore the file
            if contains is not None and filename.find(contains) == -1:
                continue

            # determine the file extension of the current file
            ext = filename[filename.rfind("."):].lower()

            # check to see if the file is an image and should be processed
            if ext.endswith(validExts):
                # construct the path to the image and yield it
                imagePath = os.path.join(rootDir, filename).replace(" ", "\\ ")
                yield imagePath

def load_images(directory='', size=(64,64)):
    images = []
    labels = [] # Integers corresponding to the categories in alphabetical order
    label = 0

    imagePaths = list(list_images(directory))

    for path in imagePaths:
        if not('OSX' in path):
            path = path.replace('\\', '/')

            image = cv2.imread(path) #Reading the image with OpenCV
            image = cv2.resize(image, size) #Resizing the image, in case some are not of the same size

            images.append(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))

    return images

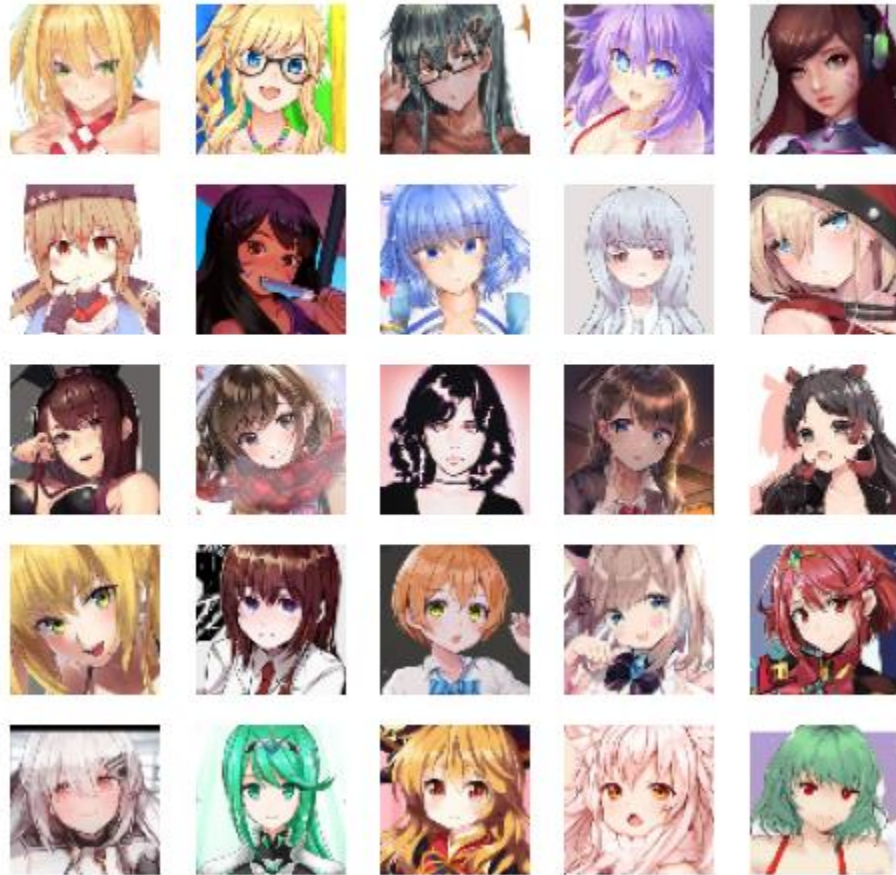
!pip install opendatasets --upgrade --quiet
import opendatasets as od

dataset_url = 'https://www.kaggle.com/datasets/scribbless/another-anime-face-dataset'
od.download(dataset_url) # {"username": "jjggggg", "key": "dd2bf9a60631eb54931473f3e45347b4"}

[ ] images = load_images('/content/another-anime-face-dataset')
```

3. We then set a dimension of 5x5 to display 25 images in the dataset.

```
[ ] _,ax = plt.subplots(5,5, figsize = (8,8))
    for i in range(5):
        for j in range(5):
            ax[i,j].imshow(images[5*i+j])
            ax[i,j].axis('off')
```



4. We created a function for building our GAN model. Here, we build the generator network as well as the discriminator network. The generator network is the one responsible for generating the fake data/images while the discriminator acts as a classifier to identify or distinguish the real data part of the generated image.

```
[ ] class GAN():
    def __init__(self):
        self.img_shape = (64, 64, 3)

        self.noise_size = 100

        optimizer = Adam(0.0002, 0.5)

        self.discriminator = self.build_discriminator()
        self.discriminator.compile(loss='binary_crossentropy',
                                    optimizer=optimizer,
                                    metrics=['accuracy'])

        self.generator = self.build_generator()
        self.generator.compile(loss='binary_crossentropy', optimizer=optimizer)

        self.combined = Sequential()
        self.combined.add(self.generator)
        self.combined.add(self.discriminator)

        self.discriminator.trainable = False

        self.combined.compile(loss='binary_crossentropy', optimizer=optimizer)

        self.combined.summary()

# Creating the generator, the large kernels in the convolutional layers allow the network to create complex structures.
def build_generator(self):
    epsilon = 0.00001 # Small float added to variance to avoid dividing by zero in the BatchNorm layers.
    noise_shape = (self.noise_size,)

    model = Sequential()

    model.add(Dense(4*4*512, activation='linear', input_shape=noise_shape))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Reshape((4, 4, 512)))

    model.add(Conv2DTranspose(512, kernel_size=[4,4], strides=[2,2], padding="same",
                               kernel_initializer=keras.initializers.TruncatedNormal(stddev=0.02)))
    model.add(BatchNormalization(momentum=0.9, epsilon=epsilon))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2DTranspose(256, kernel_size=[4,4], strides=[2,2], padding="same",
                               kernel_initializer=keras.initializers.TruncatedNormal(stddev=0.02)))
    model.add(BatchNormalization(momentum=0.9, epsilon=epsilon))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2DTranspose(128, kernel_size=[4,4], strides=[2,2], padding="same",
                               kernel_initializer=keras.initializers.TruncatedNormal(stddev=0.02)))
    model.add(BatchNormalization(momentum=0.9, epsilon=epsilon))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2DTranspose(64, kernel_size=[4,4], strides=[2,2], padding="same",
                               kernel_initializer=keras.initializers.TruncatedNormal(stddev=0.02)))
    model.add(BatchNormalization(momentum=0.9, epsilon=epsilon))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2DTranspose(3, kernel_size=[4,4], strides=[1,1], padding="same",
                               kernel_initializer=keras.initializers.TruncatedNormal(stddev=0.02)))
```

```

# Standard activation for the generator of a GAN
model.add(Activation("tanh"))

model.summary()

noise = Input(shape=noise_shape)
img = model(noise)

return Model(noise, img)

```

```

def build_discriminator(self):

    model = Sequential()

    model.add(Conv2D(128, (3,3), padding='same', input_shape=self.img_shape))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization())
    model.add(Conv2D(128, (3,3), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size=(3,3)))
    model.add(Dropout(0.2))

    model.add(Conv2D(128, (3,3), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization())
    model.add(Conv2D(128, (3,3), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size=(3,3)))
    model.add(Dropout(0.3))

    model.add(Flatten())
    model.add(Dense(128))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dense(128))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Dense(1, activation='sigmoid'))

    model.summary()

    img = Input(shape=self.img_shape)
    validity = model(img)

    return Model(img, validity)

```


4. We start to train the generator as well as the discriminator network for the last part of defining the GAN model. Here, the program defines the loss definition for the model.

```
def train(self, epochs, batch_size=128, metrics_update=50, save_images=100, save_model=2000):

    X_train = np.array(images)
    X_train = (X_train.astype(np.float32) - 127.5) / 127.5

    half_batch = int(batch_size / 2)

    mean_d_loss=[0,0]
    mean_g_loss=0

    for epoch in range(epochs):
        idx = np.random.randint(0, X_train.shape[0], half_batch)
        imgs = X_train[idx]

        noise = np.random.normal(0, 1, (half_batch, self.noise_size))
        gen_imgs = self.generator.predict(noise)

        # Training the discriminator

        # The loss of the discriminator is the mean of the losses while training on authentic and fake images
        d_loss = 0.5 * np.add(self.discriminator.train_on_batch(imgs, np.ones((half_batch, 1))),
                               self.discriminator.train_on_batch(gen_imgs, np.zeros((half_batch, 1))))

        # Training the generator
        for _ in range(2):
            noise = np.random.normal(0, 1, (batch_size, self.noise_size))

            valid_y = np.array([1] * batch_size)
            g_loss = self.combined.train_on_batch(noise, valid_y)

        mean_d_loss[0] += d_loss[0]
        mean_d_loss[1] += d_loss[1]
        mean_g_loss += g_loss
```



```

# We print the losses and accuracy of the networks every 200 batches mainly to make sure the accuracy of the discriminator
# is not stable at around 50% or 100% (which would mean the discriminator performs not well enough or too well)
if epoch % metrics_update == 0:
    print ("%d [Discriminator loss: %f, acc.: %.2f%%] [Generator loss: %f]" % (epoch, mean_d_loss[0]/metrics_update, 100*mean_d_loss[1]/metrics_update, mean_g_loss/metrics_update,
    mean_d_loss=[0,0]
    mean_g_loss=0

# Saving 25 images
if epoch % save_images == 0:
    self.save_images(epoch)

# We save the architecture of the model, the weights and the state of the optimizer
# This way we can restart the training exactly where we stopped
if epoch % save_model == 0:
    self.generator.save("generator_%d" % epoch)
    self.discriminator.save("discriminator_%d" % epoch)

# Saving 25 generated images to have a representation of the spectrum of images created by the generator
def save_images(self, epoch):
    noise = np.random.normal(0, 1, (25, self.noise_size))
    gen_imgs = self.generator.predict(noise)

    # Rescale from [-1,1] into [0,1]
    gen_imgs = 0.5 * gen_imgs + 0.5

    fig, axes = plt.subplots(5,5, figsize = (8,8))

    for i in range(5):
        for j in range(5):
            axes[i,j].imshow(gen_imgs[5*i+j])
            axes[i,j].axis('off')

    plt.show()

    fig.savefig("Generated_Anime/Faces_%d.png" % epoch)
    plt.close()

```

5. We started to train and let the model generate new anime faces. We set the epochs to 10,001 with images being generated and save every 1k epoch. There's a total of 11 pictures includes epochs at 1 wherein there is no anime faces generated. The generated images will then be saved in a file.

```

[7] #This folder will contain the images generated during the training
!mkdir Generated_Anime

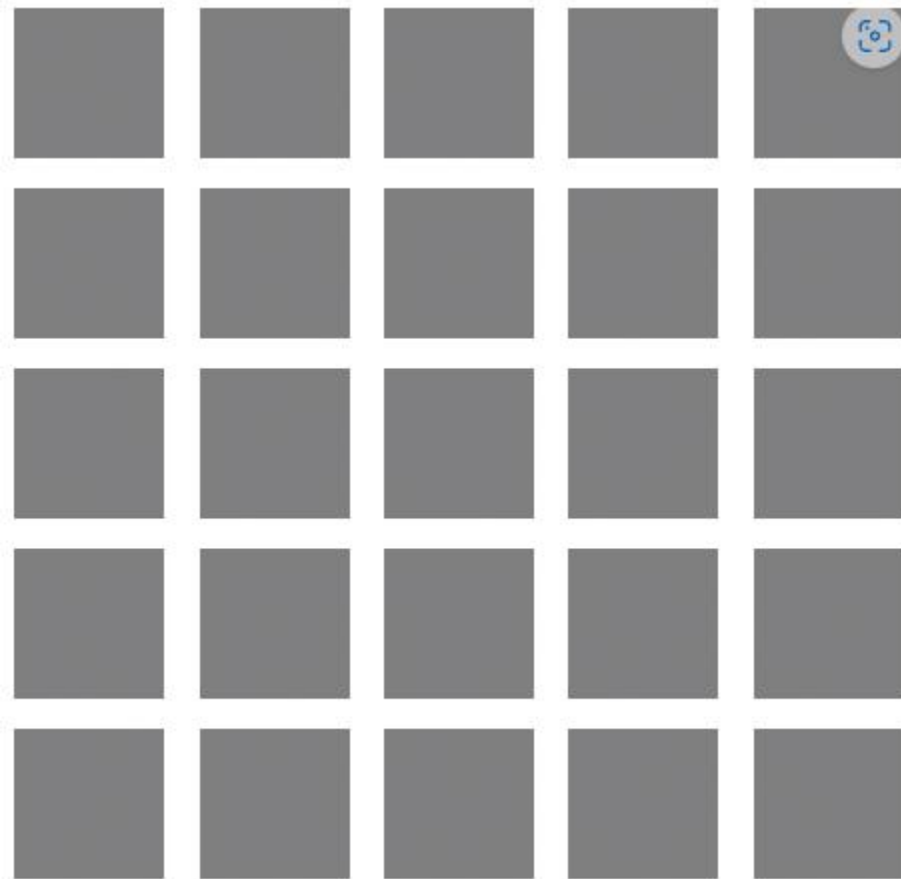
```

```

▶ gan=GAN()
gan.train(epochs=10001, batch_size=256, metrics_update=200, save_images=1000, save_model=10000)

4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 18ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 20ms/step
4/4 [=====] - 0s 19ms/step

```



Epoch at 1. No generated anime faces

1000 [Discriminator loss: 0.600680, acc.: 67.45%] [Generator loss: 1.274487
1/1 [=====] - 0s 18ms/step



Epoch at 1k. First generated anime faces

2000 [Discriminator loss: 0.556534, acc.: 71.15%] [Generator loss: 1.384334]
1/1 [=====] - 0s 15ms/step



Epoch at 2k. Generated anime faces

3000 [Discriminator loss: 0.541438, acc.: 73.72%] [Generator loss: 1.150347]
1/1 [=====] - 0s 16ms/step



Epoch at 3k. Generated anime faces

4000 [Discriminator loss: 0.577525, acc.: 69.32%] [Generator loss: 1.309778]
1/1 [=====] - 0s 22ms/step



Epoch at 4k. Generated anime faces

5000 [Discriminator loss: 0.562326, acc.: 70.94%] [Generator loss: 1.424783
1/1 [=====] - 0s 16ms/step



Epoch at 5k. Generated anime faces

6000 [Discriminator loss: 0.534986, acc.: 72.98%] [Generator loss: 1.687882]
1/1 [=====] - 0s 17ms/step



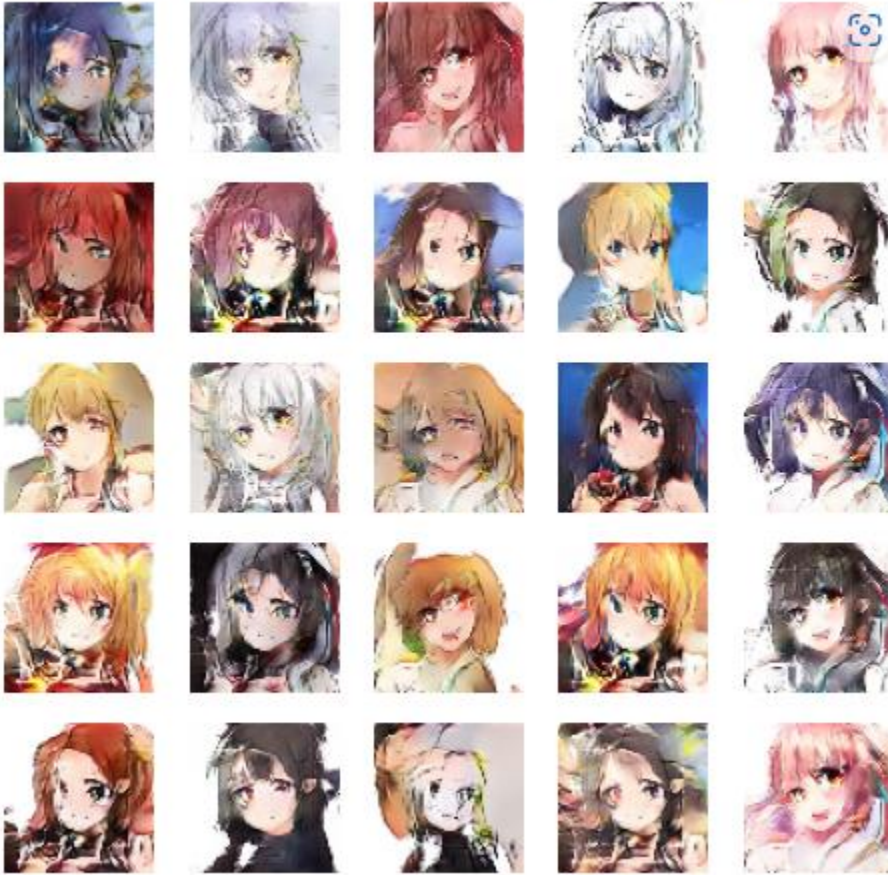
Epoch at 6k. Generated anime faces

7000 [Discriminator loss: 0.524657, acc.: 73.37%] [Generator loss: 1.781694]
1/1 [=====] - 0s 17ms/step



Epoch at 7k. Generated anime faces

```
8000 [Discriminator loss: 0.488257, acc.: 75.89%] [Generator loss: 2.071867]  
1/1 [=====] - 0s 18ms/step
```



Epoch at 8k. Generated anime faces


```
9000 [Discriminator loss: 0.481456, acc.: 76.31%] [Generator loss: 2.112323]  
1/1 [=====] - 0s 20ms/step
```



Epoch at 9k. Generated anime faces

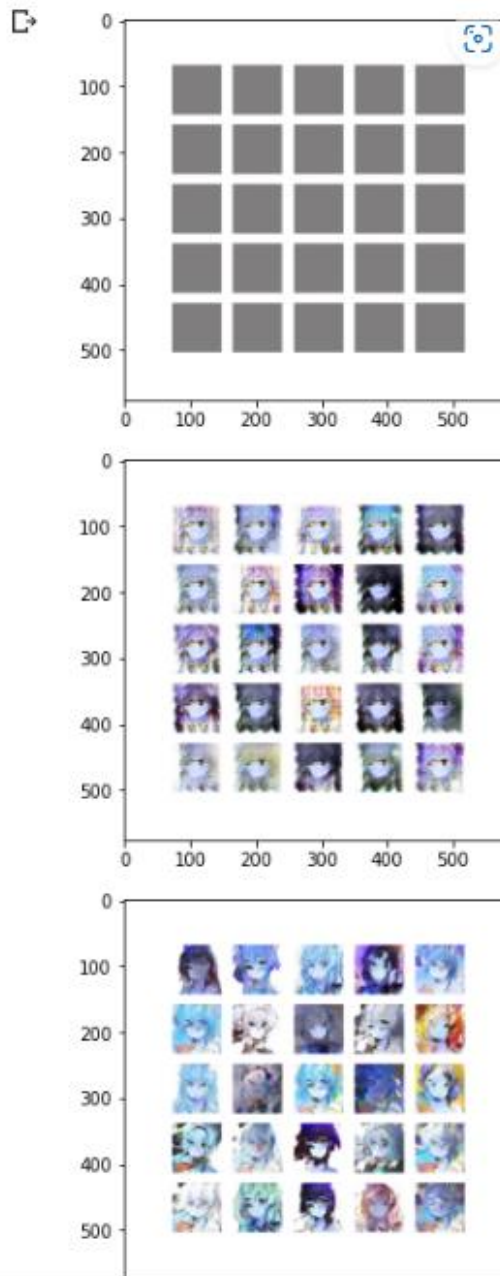
```
10000 [Discriminator loss: 0.481456, acc.: 76.31%] [Generator loss: 2.071867]  
1/1 [=====] - 0s 20ms/step
```



Epoch at 10k. Final generated anime faces for the prediction

15. We display all of the generated image and we can see that as epoch increases, the model is becoming better in generated anime faces.

```
imdir = '/content/Generated_Anime'  
ext = ['png', 'jpg', 'jpeg'] # Add image formats here  
images = os.listdir(imdir)  
  
for img in images:  
    img_arr = cv2.imread(os.path.join(imdir, img))  
    plt.figure()  
    plt.imshow(img_arr)
```



DISCUSSION:

This activity helped me to learn and understand how GAN is able to generate new examples from existing samples and one of its applications is what was done in this activity, generating anime faces. I was able to generate new anime faces just from the dataset that I have gathered. Here, we can see the step-by-step process of how the program improves its performance.

This activity also emphasizes that the use of the graphics processing unit (GPU) was essential for the accomplishment of this activity. When I first run the program without any GPU, the program took 10 minutes just to do 50 epochs while when I activated the GPU as hardware accelerator, it only took 10 minutes to do 200 epochs. Due to limited free hardware accelerator of Google Colab, I was only able to do 10k epochs but we can see that the program starts to become better at generating anime faces at this epoch. The program is able to generate anime faces with good quality. Further epochs will make the program to generate a better and more realistic anime faces.

References:

Reference for the code: <https://www.kaggle.com/code/nassimyagoub/gan-anime-faces/notebook>

Reference for the dataset: <https://www.kaggle.com/datasets/scribbless/another-anime-face-dataset>

Brownlee, J. (2019, July 12). *18 impressive applications of generative adversarial networks (Gans)*. Machine Learning Mastery. Retrieved October 21, 2022, from <https://machinelearningmastery.com/impressive-applications-of-generative-adversarial-networks/>

Ihalapathirana, A. (2022, January 13). *Generative adversarial networks for anime face generation-pytorch*. Medium. Retrieved October 21, 2022, from <https://aihalapathirana.medium.com/generative-adversarial-networks-for-anime-face-generation-pytorch-1b4037930e21>

